

Edge_Intelligence_Lab_MACSE604_

25MML0002

Adithyash BC

IN LAB 2

Part 1

To create a lightweight model(CNN/ANN) and train it on MNIST or other image datasets.

Save the model using pickle and note the weight of the model

The model chosen was CNN and the dataset chosen was mnist for the following exercise

Code:

```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.utils import to_categorical
import pickle
import numpy as np

(X_train, y_train), (X_test, y_test) = mnist.load_data()

print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")

X_train shape: (60000, 28, 28)
X_test shape: (10000, 28, 28)
```

```

X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32')
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32')
X_train = X_train / 255.0
X_test = X_test / 255.0

# One-hot encode
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)

print(f"X_train reshaped and normalized shape: {X_train.shape}")
print(f"y_train one-hot encoded shape: {y_train.shape}")

X_train reshaped and normalized shape: (60000, 28, 28, 1)
y_train one-hot encoded shape: (60000, 10)

# Build the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an `input_shape` / `input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

]: model.compile(optimizer='adam',
                 loss='categorical_crossentropy',
                 metrics=['accuracy'])

model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_3 (Conv2D)	(None, 11, 11, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_1 (Flatten)	(None, 1600)	0
dense_2 (Dense)	(None, 128)	204,928
dense_3 (Dense)	(None, 10)	1,290

Total params: 225,034 (879.04 KB)

Trainable params: 225,034 (879.04 KB)

Non-trainable params: 0 (0.00 B)

```

#training
history = model.fit(X_train, y_train, epochs=5, batch_size=32, validation_split=0.1)
# Evaluate
loss, accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f"Test Loss: {loss:.4f}")
print(f"Test Accuracy: {accuracy:.4f}")

Epoch 1/5
1688/1688 54s 31ms/step - accuracy: 0.9022 - loss: 0.3230 - val_accuracy: 0.9847 - val_loss: 0.0570
Epoch 2/5
1688/1688 51s 30ms/step - accuracy: 0.9852 - loss: 0.0499 - val_accuracy: 0.9858 - val_loss: 0.0502
Epoch 3/5
1688/1688 82s 30ms/step - accuracy: 0.9906 - loss: 0.0313 - val_accuracy: 0.9903 - val_loss: 0.0361
Epoch 4/5
1688/1688 51s 30ms/step - accuracy: 0.9929 - loss: 0.0226 - val_accuracy: 0.9900 - val_loss: 0.0366
Epoch 5/5
1688/1688 50s 29ms/step - accuracy: 0.9950 - loss: 0.0157 - val_accuracy: 0.9895 - val_loss: 0.0368
Test Loss: 0.0360
Test Accuracy: 0.9890

model_filename = 'mnist_cnn.pkl'
with open(model_filename, 'wb') as file:
    pickle.dump(model, file)

print(f"Model saved to {model_filename} using pickle.")

Model saved to mnist_cnn.pkl using pickle.

```

The weight/ size of the model can be found by navigating to the desired model via the Jupyter notebook interface

[mnist_cnn.pkl](#)

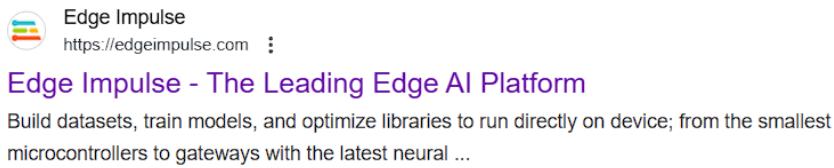
3 minutes ago 2.74 MB

The size of the model is 2.74MB

Part 2

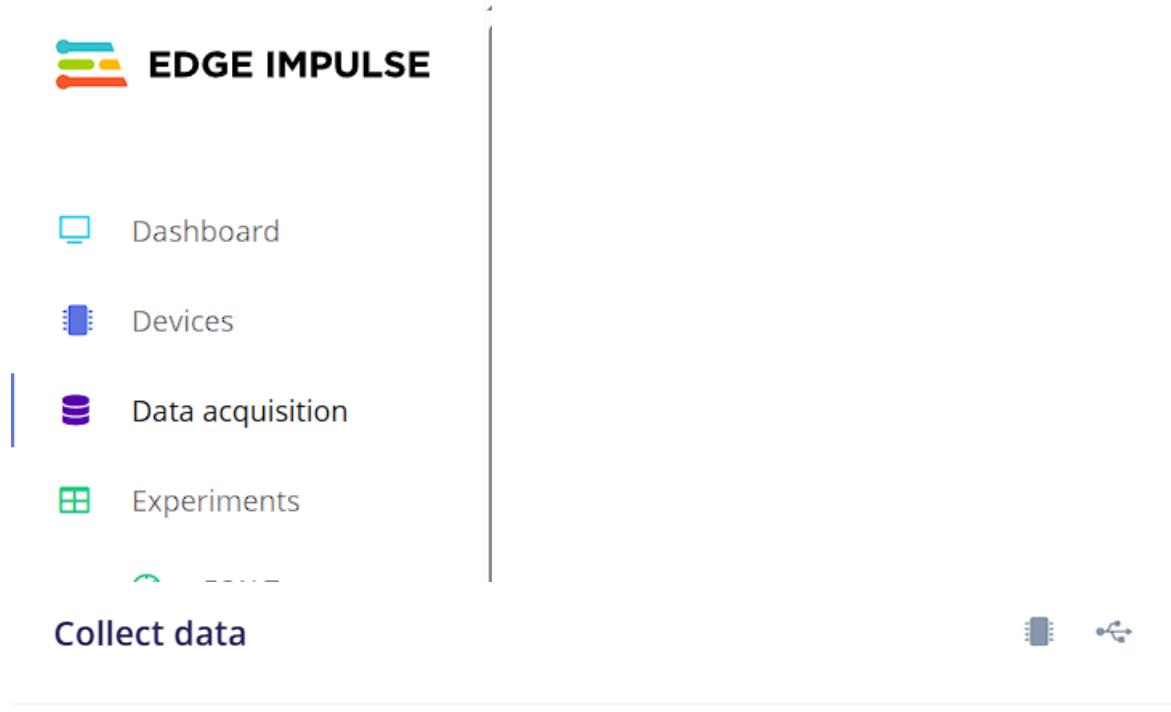
Create an edge impulse account and create an image dataset by capturing images via your cellphone

We must first navigate to edgeimpulse.com and sign up for an account



The screenshot shows the Edge Impulse website homepage. At the top, there is a navigation bar with a logo, the text "Edge Impulse", the URL "https://edgeimpulse.com", and a menu icon. Below the navigation bar, the title "Edge Impulse - The Leading Edge AI Platform" is displayed in a large, bold, purple font. A subtitle "Build datasets, train models, and optimize libraries to run directly on device; from the smallest microcontrollers to gateways with the latest neural ..." follows in a smaller, gray font.

After filling in the required details we navigate to the data acquisition tab and click on “Connect a Device” which will prompt us with three options.



The screenshot shows the Edge Impulse web interface. On the left, there is a vertical sidebar with four main options: "Dashboard", "Devices", "Data acquisition" (which is highlighted with a blue vertical bar), and "Experiments". Below the sidebar, the main content area has a title "Collect data" and a call-to-action button "Connect a device".

 [Connect a device](#) to start building your dataset.

We pick the QR code option/ Mobile and scan the displayed QR code to connect our mobile device to the platform via an API

Collect new data

X

Collect data directly from your phone, computer, device, or development board.



Scan QR code to connect to your phone

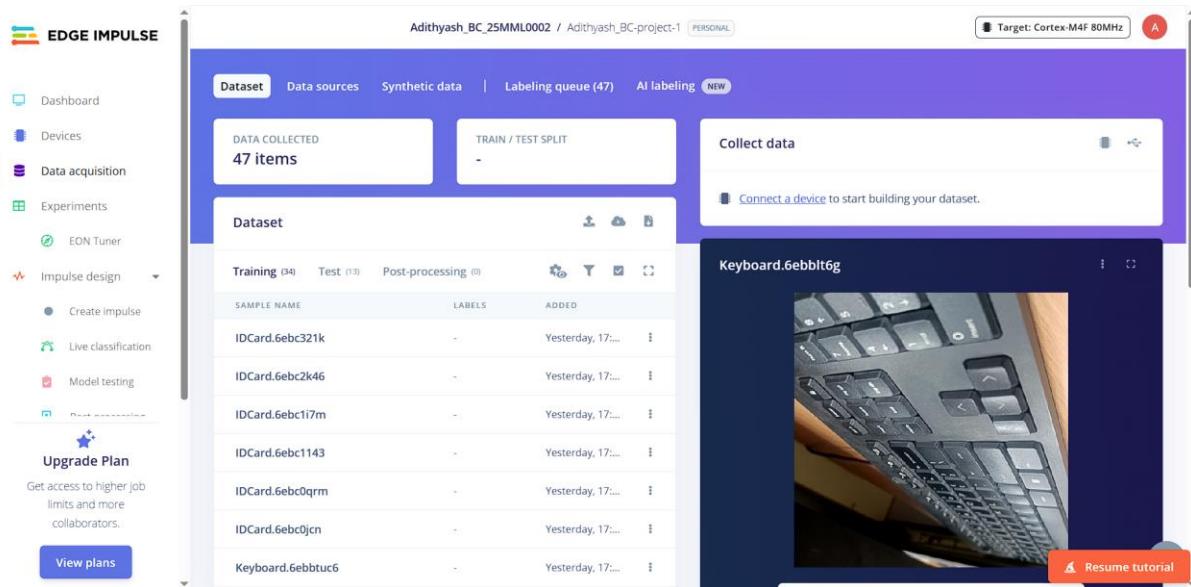


Connect to your computer



Connect your device or development board

After that we may freely take photos of objects we want in our dataset and label them. The final view should look a little something like this



The screenshot shows the Edge Impulse web interface. On the left, there's a sidebar with various options: Dashboard, Devices, Data acquisition (which is selected), Experiments, EON Tuner, Impulse design (with sub-options Create impulse, Live classification, Model testing), and an Upgrade Plan section. The main area is titled "Dataset" and shows a summary: "DATA COLLECTED 47 items" and "TRAIN / TEST SPLIT". Below this is a table for the "Dataset" tab, showing columns for "Training" (34), "Test" (13), "Post-processing" (0), and sample details like "SAMPLE NAME", "LABELS", and "ADDED". The table lists several samples, all added yesterday at 17:00. To the right, there's a "Collect data" section with a camera icon and the text "Connect a device to start building your dataset." Below this is a preview window titled "Keyboard.6ebbit6" showing a close-up photo of a keyboard. At the bottom right of the preview window is a red button labeled "Resume tutorial".

Edge In-Lab part 2

Adithyash BC

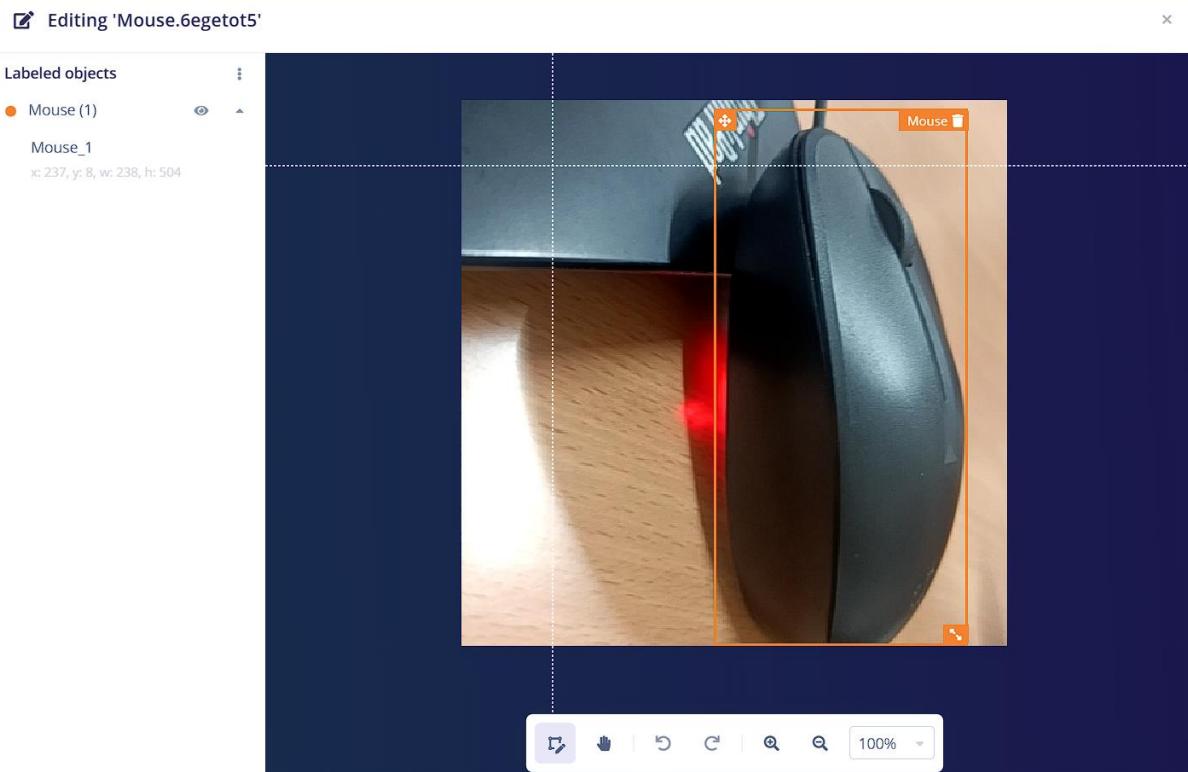
25MML0002

After the data has been acquired, it must be labelled. This can be done by navigating to the 3 dots near the images and selecting it and clicking on edit labels.

The screenshot shows the Edge Impulse web interface. On the left is a sidebar with various navigation options: Dashboard, Devices, Data acquisition, Experiments, EON Tuner, Impulse design (with sub-options Create impulse, Image, Object detection), and Help. The main area is titled "Dataset" and contains a table with three columns: SAMPLE NAME, LABELS, and ADDED. The table has seven rows, each representing a sample with its label and addition date. A context menu is open over the last row (IDCard.6ebc321k, ID, Today, 16:52:22). The menu options are: Rename, Edit labels (highlighted in blue), Clear labels, Move to test set, Move to post-processing set, Disable, Download, and Delete.

SAMPLE NAME	LABELS	ADDED
Mouse.6egetot5	Mouse	Today, 16:52:22
Mouse.6egetb82	Mouse	
Mouse.6gegesjr4	Mouse	
Mouse.6egese3q	Mouse	
Mouse.6egertug	Mouse	
Fan.6gegeoeco	Fan	
IDCard.6ebc321k	ID	

A window will pop up of the image. You can use the “plus” shaped crosser to box out the object to be labelled. Once boxed a pop up window will show up allowing you to change/make the label. This will help to better inform the model of what it is looking for.



Repeat this process with all the images in the dataset

You will know if the label has been placed by looking at the “label” column in the dataset view

LABELS
Mouse
Fan
ID

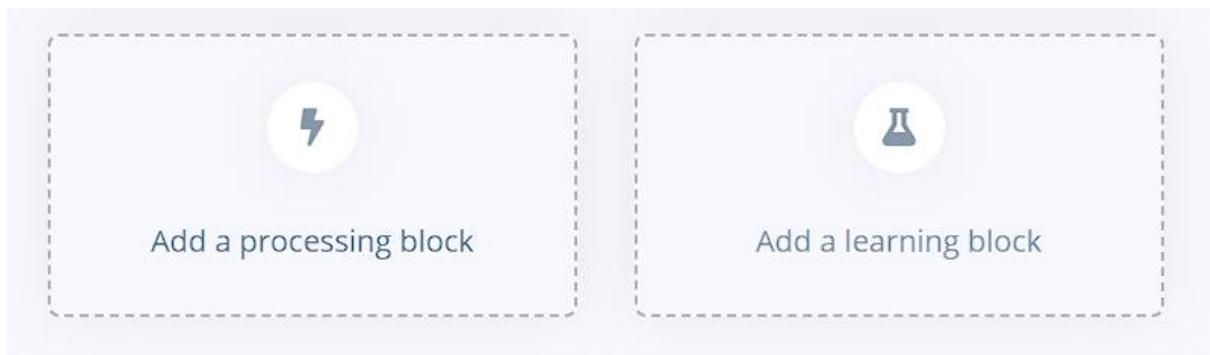
Ensure that the dataset is split into training and testing. The provided infographics on top of the dataset should be able to guide you into better managing this split



Once the data has been optimally split, navigate to the impulse tab and select on create an impulse

- Devices
- Data acquisition
- Experiments
- EON Tuner
- Impulse design ▼
 - Create impulse
 - Image Create impulse
 - Object detection
- Retrain model

Click the “Add a processing block” button to select a format the model can process. Once that is done, click the “Add a learning block” button to select a model



⚡ Add a processing block ×

Did you know? You can bring your own DSP code.

DESCRIPTION	AUTHOR	RECOMMENDED
Image <small>OFFICIALLY SUPPORTED</small> Preprocess and normalize image data, and optionally reduce the color depth.	Edge Impulse	★ Add
EEG <small>OFFICIALLY SUPPORTED</small> Filters noise and extracts spectral power features from EEG signals.	Edge Impulse	Add

Some processing blocks have been hidden based on the data in your project. [Show all blocks anyway](#)

⚡ Add a custom block Cancel

⚠ Add a learning block ×

Did you know? You can bring your own model in PyTorch, Keras or scikit-learn.

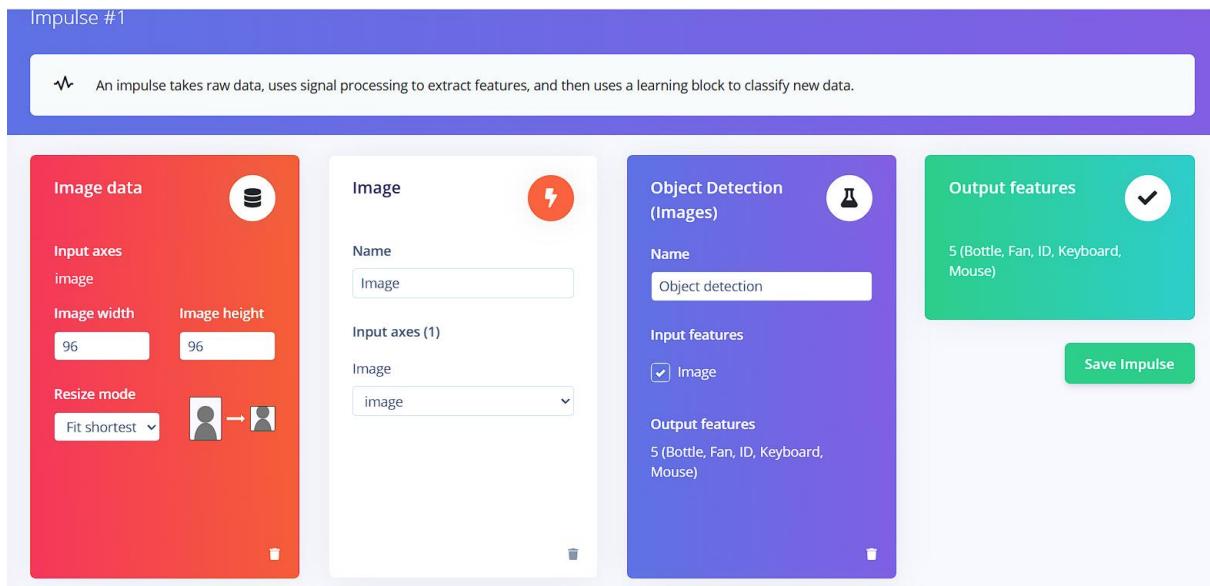
DESCRIPTION	AUTHOR	RECOMMENDED
Object Detection (Images) <small>OFFICIALLY SUPPORTED</small> Fine tune a pre-trained object detection model on your data. Good performance even with relatively small image datasets.	Edge Impulse	Add
Visual Anomaly Detection - FOMO-AD <small>OFFICIALLY SUPPORTED</small> Detect visual anomalies. Extracts visual features using a pre-trained backbone, and applies a scoring function to evaluate how anomalous a sample is by comparing the extracted features to the learned model. Does not require anomalous data.	Edge Impulse	Add

 Want access to all learning blocks? [Upgrade now.](#)

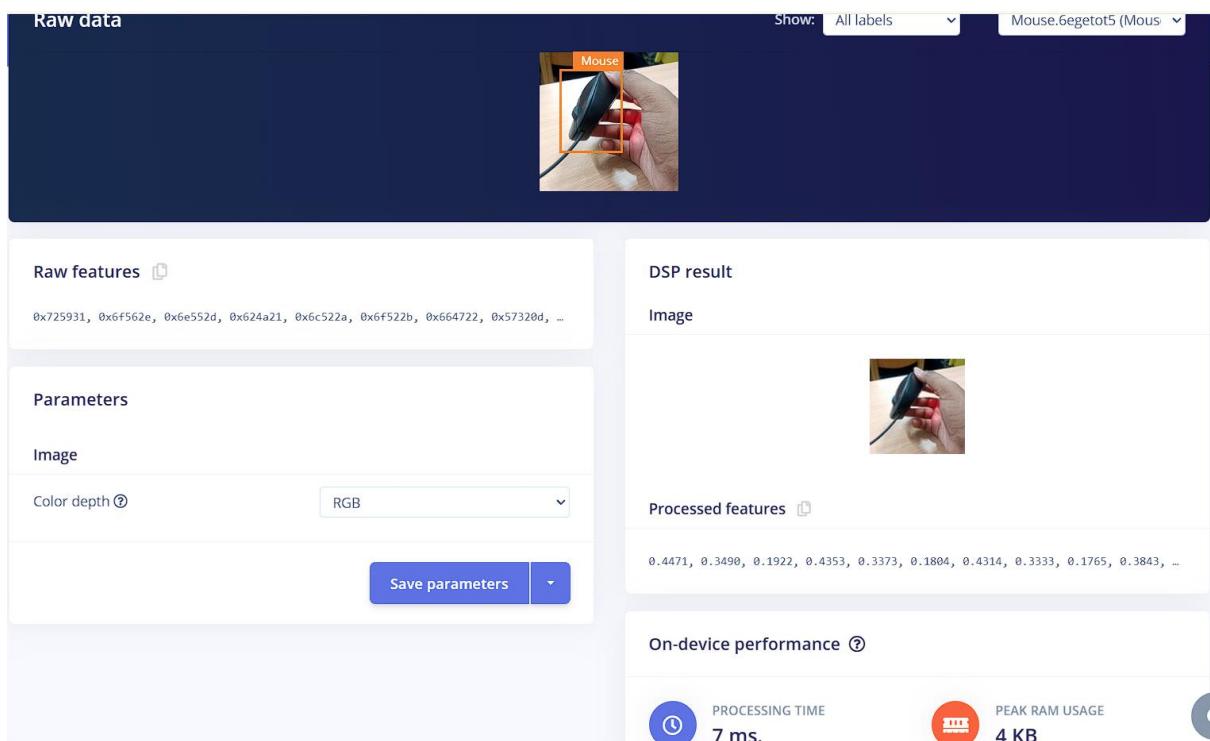
Some learning blocks have been hidden based on the data in your project. [Show all blocks anyway](#)

Once that is done, click on “the Save impulse” button to save the setup

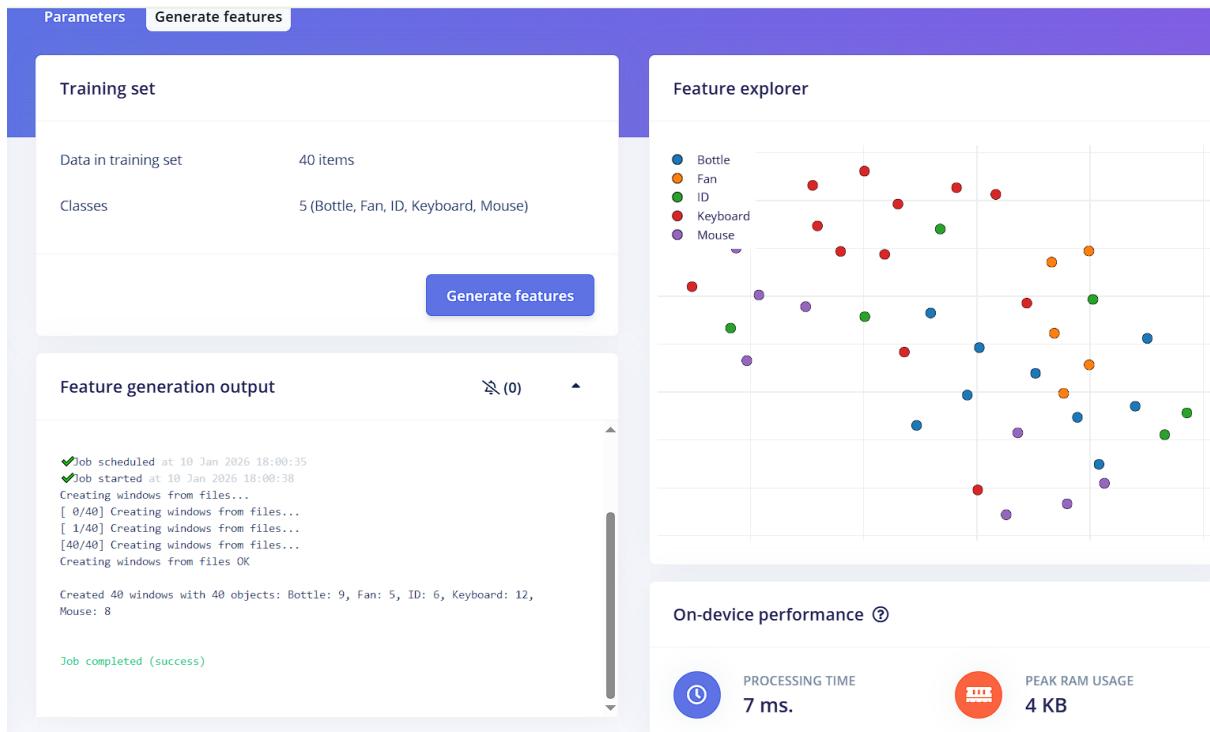
The final view should look like this



Navigate to the Image tab(below the impulse) and then save the parameters for the raw data. The parameters can be to process the image in RGB or in grayscale. Click on “save parameters” before you proceed any further.



The tab will automatically navigate to the generate features tab. Once there you must click on generate features to load the features of your dataset. After it is done it should look a little something like this



We can then navigate to the final tab, the “Object Detection” tab to run the model on our dataset.

There are various settings that can be found in the neural networks column that will allow you, the end user, to fine tune the model to your liking. Once all the nuances are dialled in(in this case we are just running it plain with no added changes), click on the “Save and train” button by navigating to the bottom of the page.

Neural Network settings

Training settings

Number of training cycles 60

Use learned optimizer

Learning rate 0.001

Training processor CPU

Data augmentation

Advanced training settings

Neural network architecture

Input layer (27,648 features)

 FOMO (Faster Objects, More Objects) MobileNetV2 0.35

Choose a different model

Output layer (5 classes)

Save & train

After a few minutes the Model should spit out the outputs, which can be found on the right-hand side of the page(under training outputs). Take note of the results and make corrections to the model or dataset if need be.

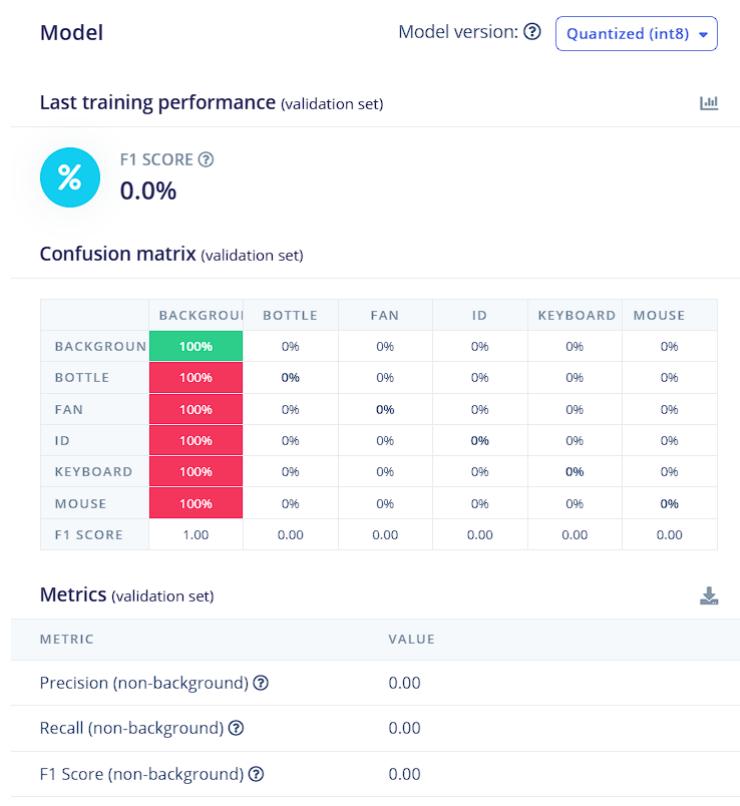
Training output

 (0)

```
INFO: Created TensorFlow Lite XNNPACK delegate for CPU.  
Calculating inferencing time OK  
Calculating float32 accuracy...  
INFO: Created TensorFlow Lite XNNPACK delegate for CPU.  
Calculating inferencing time OK  
Calculating float32 accuracy...  
INFO: Created TensorFlow Lite XNNPACK delegate for CPU.  
Calculating int8 accuracy...  
Extracting TensorBoard logs...  
Extracting TensorBoard logs OK
```

Model training complete

Job completed (success)



In our case the model gave us an F1 score of 0, which means that we will have to tune our dataset and model again to get our desired outputs.

The main reason why we were getting a low recall and f1 score was because our dataset was too small and thus the model did not have more to learn from. An easy fix to see if the model is working is to just increase the dataset size of one class. In this case I chose to increase the number of bottles

Labels in your dataset ②



This small change gave us better results and showed me that we were on the right track as right after this we got this result



F1 SCORE

16.7%

Confusion matrix (validation set)

	BACKGROUND	BOTTLE	FAN	ID	KEYBOARD	MOUSE
BACKGROUND	100%	0%	0%	0%	0%	0%
BOTTLE	83.3%	16.7%	0%	0%	0%	0%
FAN	100%	0%	0%	0%	0%	0%
ID	100%	0%	0%	0%	0%	0%
KEYBOARD	100%	0%	0%	0%	0%	0%
MOUSE	100%	0%	0%	0%	0%	0%
F1 SCORE	1.00	0.29	0.00	0.00	0.00	0.00

This confirms about our dataset theory(for now). By adding more images to the dataset



And manually annotating it, we can increase the robustness of the model and the data thus giving us better results.



We can now cycle through the same previous steps of saving impulse and generating features and then training the model to get the newer results.

Last training performance (validation set)



F1 SCORE ⓘ

45.2%

Confusion matrix (validation set)

	BACKGROUND	BOTTLE	FAN	ID	KEYBOARD	MOUSE
BACKGROUND	100%	0%	0%	0%	0%	0%
BOTTLE	100%	0%	0%	0%	0%	0%
FAN	100%	0%	0%	0%	0%	0%
ID	-	-	-	-	-	-
KEYBOARD	22.2%	0%	0%	0%	77.8%	0%
MOUSE	100%	0%	0%	0%	0%	0%
F1 SCORE	1.00	0.00	0.00		0.88	0.00

These results show us that the increased dataset pool is helping the model to better understand and recognize the images. The way to increase the F1 score is to use even more images and even better images.